

# IN-STK5000

## Project 2

Fall 2023

Helene Bøsei Olsen, Sander Finnset Ørnes,  
Even Tronstad, Andreas Christian Poole

# Agenda

- Data Leakage
- Reproducibility
- Pipeline
- Model evaluation - performance metrics
- Ranges of performance metrics
- Privacy

# Baseline model: Data leakage

We considered data leakage through all steps in our analysis:

## **Scope reduction:**

Deleted samples based on gender and race

- Based on the whole data set. **Danger!**

## **Train-Test split:**

- Done very early
- Seperate data sets through most of the process

**Data Analysis** Primarily performed only on training set

- Outliers, correlations
- For missing data performed on whole data set
  - Should not cause data leakage

# Data leakage

## Data manipulation:

- Do operations on training and test set in separate steps
- Separation of data to be corrected and data used to correct (train data)

Example:

```
train = handle_outliers(train, train_outlier_bounds)
test = handle_outliers(test, train_outlier_bounds)
```

## Feature selection:

- Mostly used non data-driven criteria:
  - Common sense and cost of collection
- Temperature removed based on low variance (train data)

# Data leakage

## Classification:

- Simple classifier with manually set hyperparameters.
  - Might have been unconsciously set based on info from test set

**Danger!**

## Correction of data leakage issues:

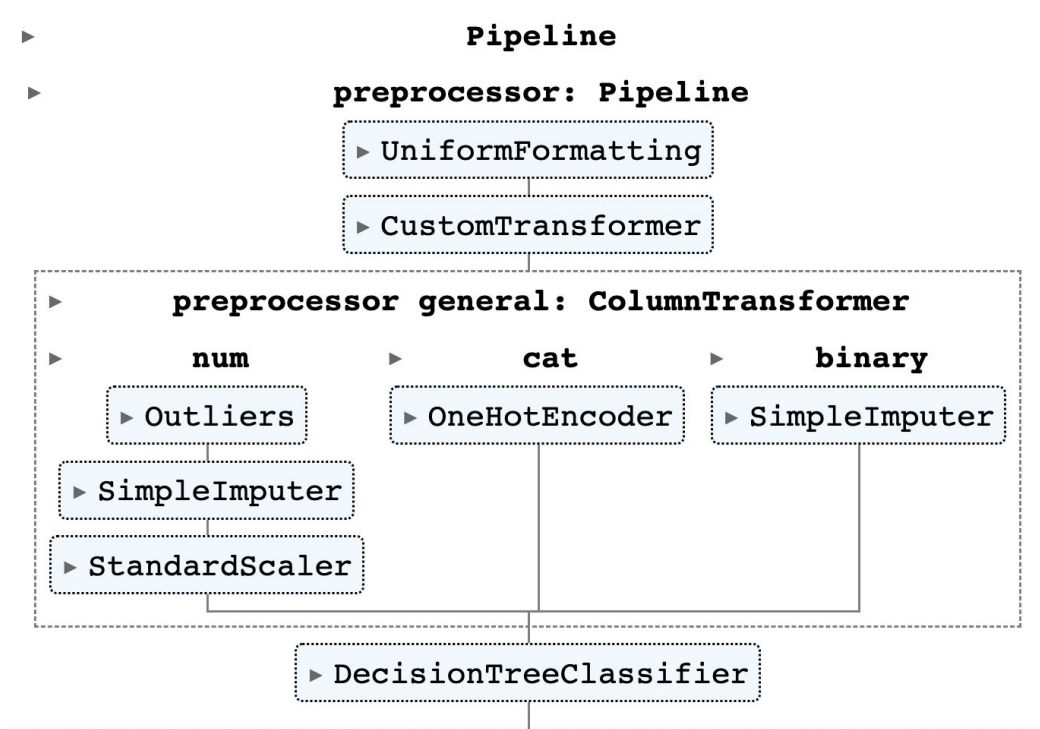
- Not delete non whites and females
- Hyperparameter tuning with grid search cross validation

# New implementation: Pipeline

- Never edits the actual data set.
  - all preprocessing and formatting in the pipeline.
- Preprocessing
  - imputes and scales without using test set by design.
- No scope reduction.
  - include gender and race (one hot encoder with 'other' if prevalence less than 0.1)
- Hyperparameter tuning with grid search cross validation.
  - on max\_depth and complexity parameter ccp\_alpha.

# Pipeline Structure

- Custom transformers imputes Obesity and Polydipsia
- General preprocessing
- Decision tree as classifier
- Finally, grid search on tree depths and complexity parameter (not displayed)



# Reproducibility

Ensure reproducibility through the following steps:

- All code and documentation available on GitHub
  - Installation script for easy usage
- Set a global seed
  - Separate seed for privacy
- requirements.txt file available with versions of packages / libraries etc.
- Thorough instructions on how to run experiments
- Simplifying code with pipelines
- Pytest for reproducibility



# Model evaluation - performance metrics

Business case recap:

- Work for the Public health authorities
- Diagnosing and treatment handled by the Private health service
- Machine Learning system on website for evaluation of diabetes risk
- Goal: Get the right people tested

Accuracy:

- Only ~10% of population have diabetes
- Can achieve high accuracy by always predicting negative
- Flawed measure of how system help us achieve our goal

# Model evaluation - performance metrics

Precision:

- High precision tells us we are not advising too many people to go into the doctor's office

Recall:

- High recall tells us we do not miss many positive cases of diabetes

F<sub>1</sub>-score:

- Balance between precision and recall
- High F<sub>1</sub> score tells us we are able to detect many cases of diabetes, without many negative tests
- Most important evaluation metric for how system help us achieve our goals

**Performance of our classifier:**

	Test set
Accuracy	92.7%
Precision	95.4%
Recall	92.5%
F1	93.9%

# Ranges and methodology

## Repeated Bootstrap

- Draw training set as bootstrap sample of the size of the entire data set
  - ~60% unique samples
- All datapoints not included in bootstrap sample is the test set
  - ~40% of total data set
- Train on training set
- Predict and evaluate performance on test set

## Why this method?

- Simple and transparent
- Can generate any number of estimates for the metrics
  - Not limited by number of folds as CV

# Results - ranges

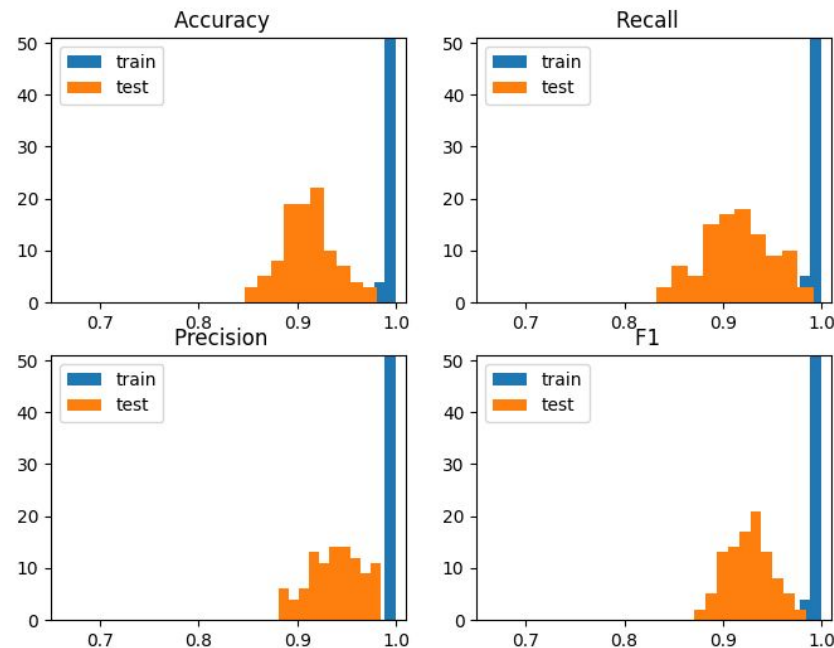
Ranges for 100 estimates of the chosen metrics

Training data almost always obtains perfect scores

Quite low spread of scores on test set  
- Low standard deviations

F1-score has the smallest standard deviation of all metrics

Model seems robust, and performs well across metrics and test sets!



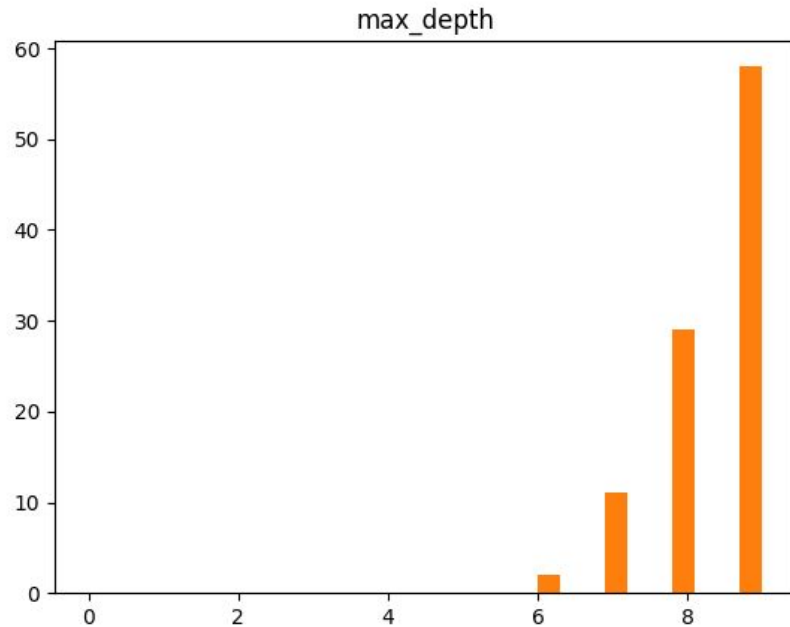
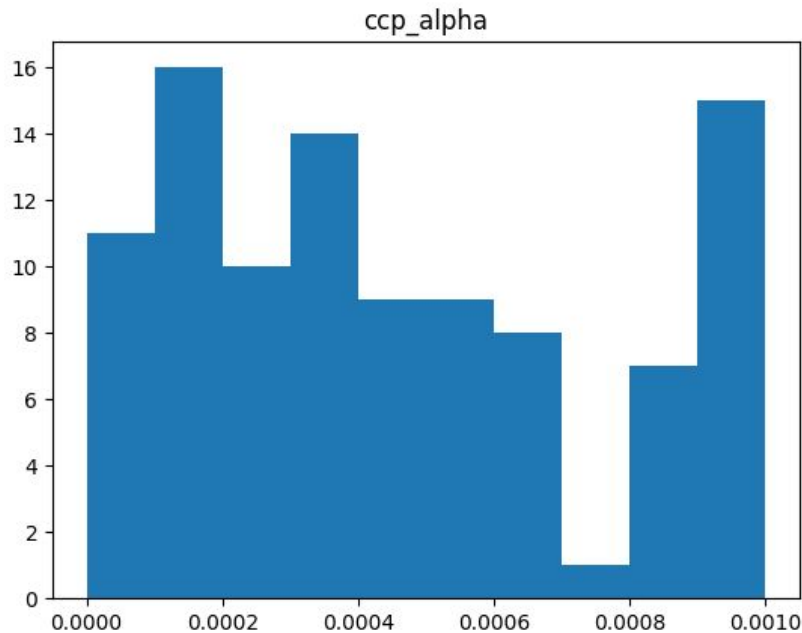
	Train mean	Train stdev	Test mean	Test stdev
Accuracy	99.7%	0.003	91.1%	0.027
Precision	99.9%	0.002	93.8%	0.026
Recall	99.6%	0.005	91.6%	0.036
F1	99.8%	0.003	92.6%	0.023

Table 1: Performance on original data

# Hyperparameters - plots

- Grid defined by tree depth and cost-complexity pruning (ccp)
- Optimal parameters found using cross-validation on grid

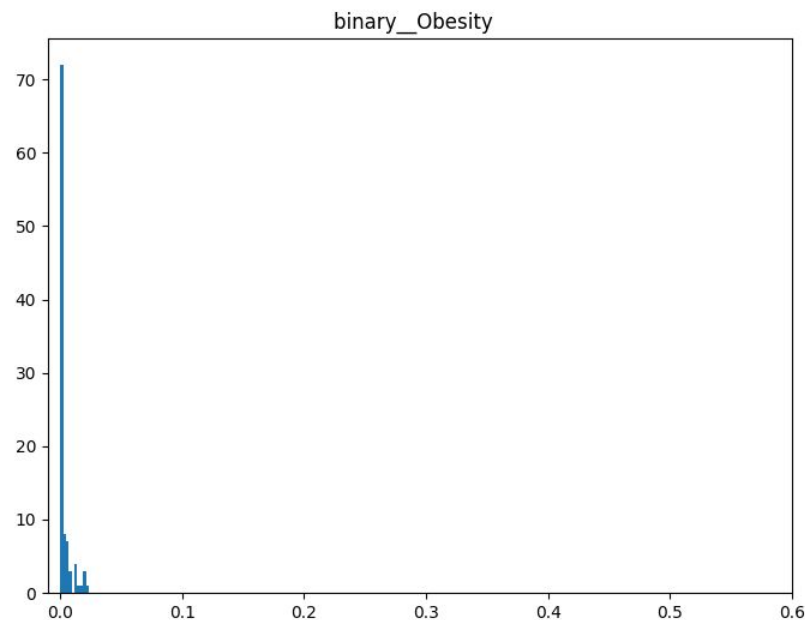
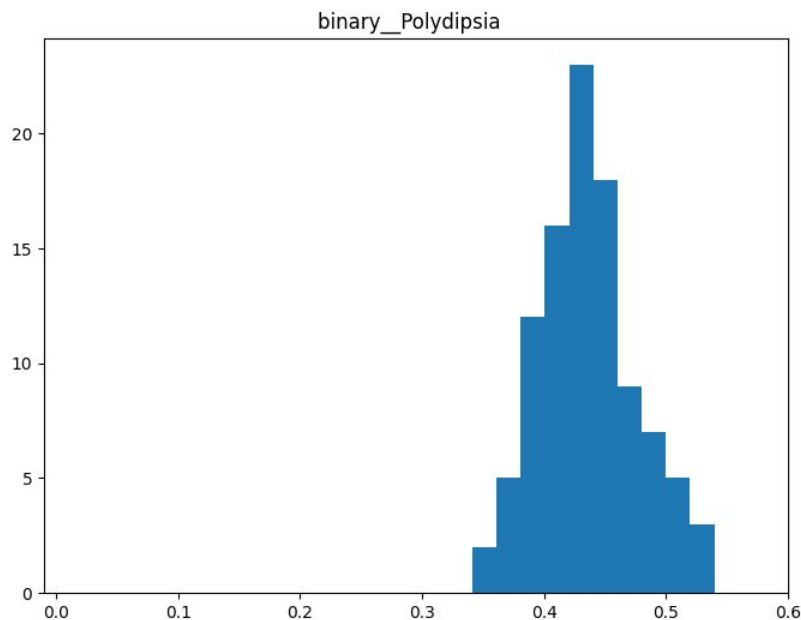
Figures show optimal parameters for each sample



# Feature importance

Built in method of sklearn

- Measures Gini importance
- Higher Gini importance means more important variable



# Privacy

Dataset: Personally identifiable health information

Anonymize the binary data (Majority of data)

- Sensitive data:
  - genital thrush, obesity, gender etc.

Adding controlled noise through Differential privacy ( $\epsilon$ -differential privacy)

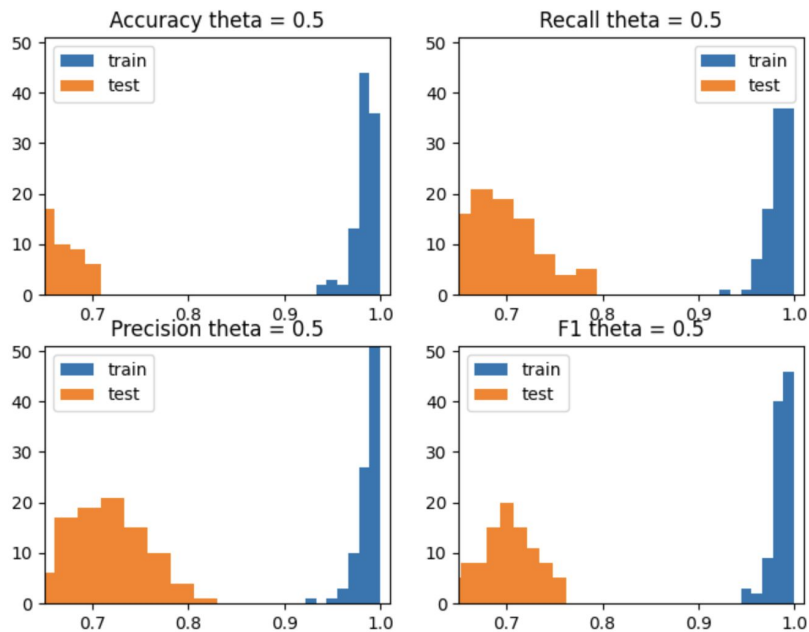
- Allow us to train a machine learning model on the data without compromising information on specific individuals
- Avoids de-anonymization with possible future datasets.
  - In contrast, K-anonymity gives no such guarantee

# Privacy procedure

- **Randomized response:**
  - Coin flipping with a probability of answering truthfully =  $\theta$
- **Column wise:**
  - Calculated independently for each column
  - measure privacy guarantee- Calculate epsilon for each column
- **Set a separate random seed for anonymization**
- **Result in new dataset - for reproducing experiments**
- **Not randomized target**
  - Small dataset
  - Less meaningful performance metrics when comparing noisy target
- **We experiment with two different cases:  $\theta = 0.5$  and  $\theta = 0.95$** 
  - $\epsilon = 1.09$  and  $\epsilon = 3.66$

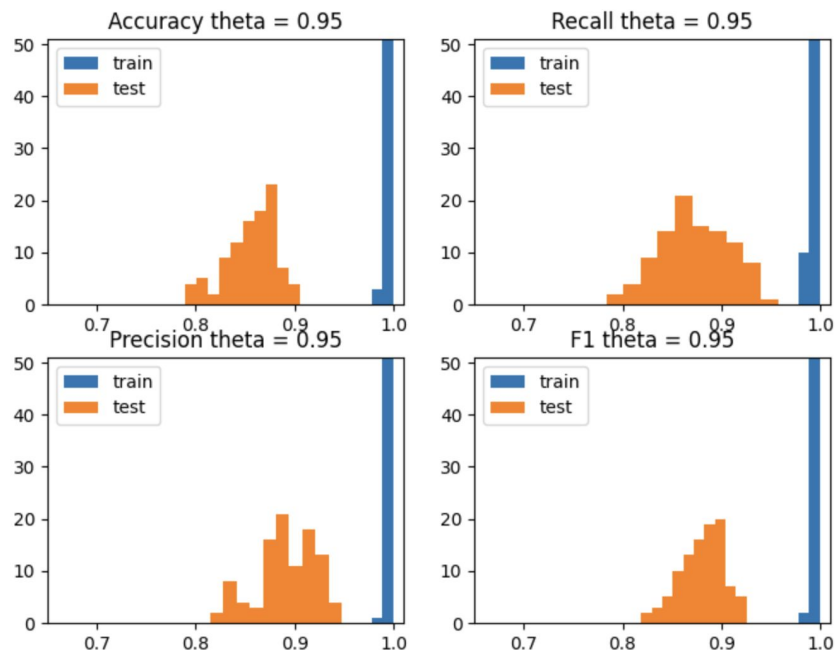


# Experiments with different theta



	Train mean	Train stdev	Test mean	Test stdev
Accuracy	98.3%	0.012	63.4%	0.037
Precision	98.9%	0.012	70.9%	0.046
Recall	98.4%	0.012	68.9%	0.044
F1	98.6%	0.010	69.7%	0.032

Table 3: Theta = 0.5



Accuracy	99.8%	0.003	<b>85.7%</b>	0.026
Precision	99.9%	0.002	<b>89.1%</b>	0.031
Recall	99.7%	0.005	<b>87.3%</b>	0.036
F1	99.8%	0.003	<b>88.1%</b>	0.022

Table 5: Theta = 0.95

# Privacy is not free

	Original data		Anonymized data	
	Test mean	Test stdev	Test mean	Test stdev
Accuracy	<b>91.1%</b>	0.027	85.7%	0.026
Precision	<b>93.8%</b>	0.026	89.1%	0.031
Recall	<b>91.6%</b>	0.036	87.3%	0.036
F1	<b>92.6%</b>	0.023	88.1%	0.022

Table 7: Comparing results on test set for original data and anonymized data with  $\theta = 0.95$

# Privacy: future work

- Differential privacy to proof against current and future datasets
  - Only applied to binary features
- Possible improvement:
  - combine differential privacy with other methods such as k-anonymity to increase the level of privacy.
  - or combine with Laplace for continuous data
  - Anonymise target